

Louis Ledoux

Philosophiae Doctor

Barcelona, 08014

Spain

* 26 February 1995

☎ +33 7 70 49 11 98

✉ i.f.ledoux@gmail.com

🌐 bynaryman.github.io


in ledoux-louis

🐦 @L0u1s73doux

🔗 Bynaryman

ID 0009-0006-3643-2157

🔗 JWBIqG8AAAAJ



“Post Hoc Ergo Propter Hoc”

Research Interests

Computer Architecture: Floating-Point Units, Systolic Arrays, Matrix-Matrix Multiply (MMM) Units, GPUs, FPGAs

Computer Arithmetic: Number Representations, Application-Specific Circuits, FloPoCo, Posit & IEEE754 Standards, Kulisch Accumulator, Accuracy and Energy Budgeting

High Performance Computing: BLAS, GEMMs, Heterogeneous Workloads, Numerical Analysis, Supercomputing

Education

2018 – 2024 **PhD Student**, *Universitat Politècnica de Catalunya (UPC)*, Barcelona

Supervisor: Marc Casas Guix. Thesis title: “Floating-Point Arithmetic Paradigms for High-Performance Computing: Software Algorithms and Hardware Designs”. Jury members attributed the **excellent** grade.

2015 – 2018 **Engineering School**, *ESIR*, Université de Rennes

Completed a three-year program culminating in an **Engineer diploma** certified by the CTI (Comité des Titres d’Ingénieurs) and a **Master’s degree** (Magister) in Computer Science.

2013 – 2015 **Classe Préparatoire**, *CUPGE ESIR*, Université de Rennes

Intensive program with a strong emphasis on Mathematics and Computer Sciences.

Experience

2018 – Now **Researcher**, *Barcelona Supercomputing Center (BSC) - RoMoL/CAOS/SONAR*, Barcelona

During the completion of my PhD, I also worked as a researcher, publishing peer-reviewed papers and attending international conferences. The conducted research is summarized as follows:

- Explored co-designed hardware/software acceleration of posit arithmetic for FPGA and ASIC technologies, and Power9 host.
- Developed Kulisch/Quire accumulators for any floating-point representation.
- Conducted accuracy and energy budgeting tailored to workload numerical requirements.
- Designed Systolic Array architecture for HPC workloads with three directions of data flow, infinite number representations, and tailored internal precision.
- Proposed very slow and small division units targeting the SIMD/Vector paradigm at the architecture level, further improving the parallelism/latency/energy ratio.

2017 - 2018 **Hardware Engineer**, *b\com*, Rennes

Coop student under a professionalization contract during my final year of study. Engaged in one year of R&D focused on FPGA acceleration in the cloud, aiming to evaluate the feasibility of these novel solutions.

- Successfully integrated an IP (real-time SDR to HDR) to convert video from NVMe to NVMe through cloud FPGA.
- Developed the IP integration using HDLs and tweaked PCI-e drivers (EDMA and XDMA).
- Scheduled user processes asynchronously with OpenCL to overlap reads and writes, successfully saturating PCIe bandwidth (≈ 15.8 GB/s).
- Traveled to Xilinx XDF Frankfurt 2018 to conduct operational monitoring and analysis.

Keywords: SDAccel, OpenCL, Xilinx FPGA (Ultrascale VU9P), VHDL, SystemVerilog, AWS F1 instances, Linux driver, Linux kernel, PCIe, C++1x, C, DMA, AaaS (Acceleration as a Service), FaaS (FPGA as a Service), Virtual Machine, Docker, PCI-e virtual functions.

July 2017 **Back End Developer**, *WaryMe*, Rennes

Summer internship focused on developing the entire back end of a people security application.

- Designed and implemented backend services and APIs.
- Ensured secure data transmission using HTTPS and Let’s Encrypt.
- Deployed and managed the application on AWS with PM2 and Nginx.
- Automated deployment processes with Jenkins.
- Collaborated with front-end developers using Angular and TypeScript.

Keywords: Node.js, Git, C++, C, Angular, TypeScript, HTTPS, Let’s Encrypt, PM2, Nginx, AWS, Jenkins.

July 2016 **Back End Developer**, ASKIA, Paris 10ème

During this summer internship, I developed an automated CLI tool for publishing surveys on popular platforms such as GitHub and Zendesk. **Key Concepts:** Node.js, HTTP(S), REST API, Git, Test-Driven Development (TDD), Event-Driven/Asynchronous Programming, Jasmine Framework, Mocks, Stubs.

July 2014 **Electronics Technician**, *Radio Electronique Rennaise (R.E.R)*, Rennes

As an electronics technician, I was responsible for repairing various electronic devices from customers, with an emphasis on audio equipment. Tasks included soldering, reverse engineering amplifier circuits, welcoming customers, and sorting/ordering electronic parts.

Peer-reviewed Conference Papers

[LC23a] **L. Ledoux** and M. Casas, “An Open-Source Framework for Efficient Numerically-Tailored Computations,” in *2023 33rd International Conference on Field-Programmable Logic and Applications (FPL)*, Gothenburg, Sweden. doi: 10.1109/FPL60245.2023.00011, arXiv:2406.02579, HAL:04277512

[LC22] **L. Ledoux** and M. Casas, “A Generator of Numerically-Tailored and High-Throughput Accelerators for Batched GEMMs,” in *2022 IEEE 30th Annual International Symposium on Field-Programmable Custom Computing Machines (FCCM)*, May 2022, pp. 1–10, New York, USA. doi: 10.1109/FCCM53951.2022.9786164, HAL:04103774

Poster Presentations

[LC24b] **L. Ledoux** and M. Casas, “LLMMMM: Large Language Models Matrix-Matrix Multiplications Characterization on Open Silicon,” in *2024 11th BSC Symposium*, Barcelona, Spain. HAL:04592229



[LC24a] **L. Ledoux** and M. Casas, “The Grafted Superset Approach: Bridging Python to Silicon with Asynchronous Compilation and Beyond,” in *2024 4th Workshop on Open-Source Design Automation (OSDA)*, hosted at DATE, Valencia, Spain. HAL:04587458


[LC23b] **L. Ledoux** and M. Casas, “Open-Source GEMM Hardware Kernels Generator: Toward Numerically-Tailored Computations,” in *2023 10th BSC Symposium*, Barcelona, Spain. arXiv:2305.18328, HAL:04094835

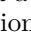
Invited Talks


[LC19] **L. Ledoux** and M. Casas, “Accelerating DL Inference with (Open)CAPI and Posit Numbers,” in *OpenPOWER Summit 2019*, Lyon, France. HAL:04094850


Projects & Hobbies

MPW 5 Taped out a Systolic Array for Matrix Multiplication with *Posit* numbers and *Quires* accumulators . I have been interviewed on YouTube .


MPW 1 My first taped-out chip on the first-ever open shared Multi-Project Wafer in collaboration with Google, SkyWater, and Efabless. Open contribution to the Open Source Silicon community available on GitHub .

SUF Developed a Python-to-ASIC compiler with a focus on arithmetic. Includes novel (and artistic) placement visualization , among other features such as plot-automation, PDK-agnosticism, asynchronous and parallel execution of several designs, etc.

OSFNTC Developed an Open-Source Framework for Efficient Numerically-Tailored Computation. Modified PyTorch, Numpy, OpenBLAS, OpenCAPI to work with PCI-e/NVLink FPGAs with arbitrarily large Systolic Arrays and infinite float format representations with any intermediate precision. Systematically improved energy efficiency and accuracy for drastically different workloads such as Deep Learning and Sea Surface Height calculation. Project on GitHub .

POF Designed the Posit Operators Framework, a comprehensive SW/HW co-designed library to facilitate arithmetic computations and neural network operations using the *Posit* numerical format on FPGAs. Manually written modules can be found on GitHub .

VH2V Designed a VHDL-to-Verilog translation tool tailored to convert FloPoCo outputs to OpenLane/OpenROAD inputs. Project available on GitHub .

Synthesizers Crafted analog and digital synthesizers for modular synthesis and Eurorack systems. Responsibilities included PCB manufacturing , part ordering, and digital design.

Languages

French Native

Spanish Native (with an honest French accent)

English Full Proficiency